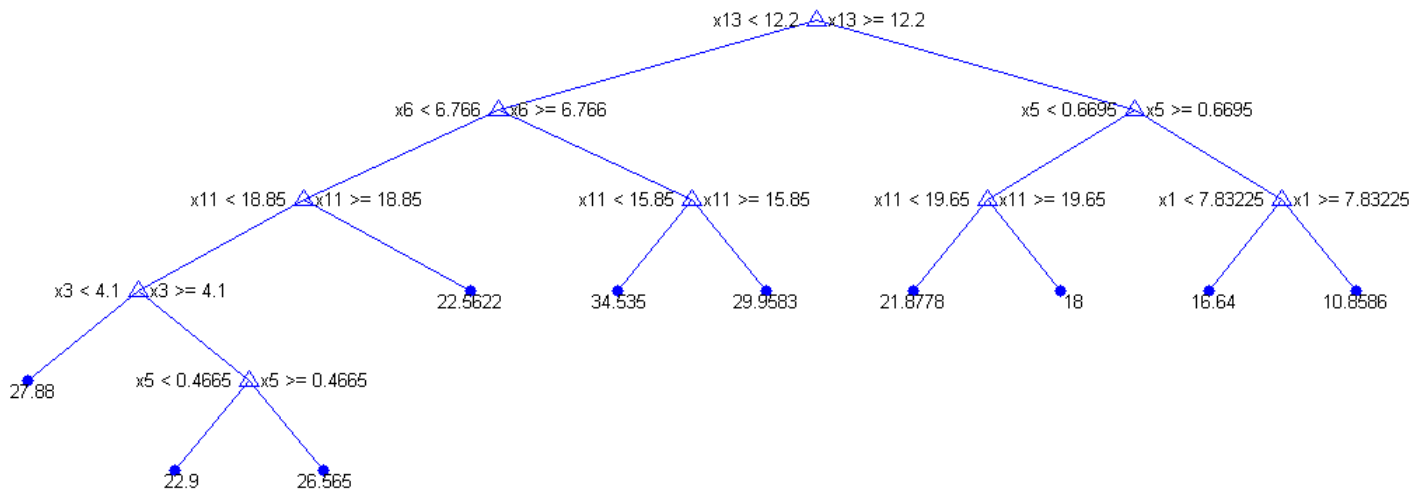

Matlab 3

Mikhail Andreev

November 10, 2015

1 EXPLORING BOSTON HOUSING DATA WITH REGRESSION TREES

The regression tree generated from the training data can be seen here:




Using the input vector:

[5, 18, 2.31, 1, 0.5440, 2, 64, 3.7, 1, 300, 15, 390, 10]

We get an output $MEDV = 22.047619$.

If we plot the Mean Absolute Error for both the training and testing sets, using different numbers of observations per leaf, you get the following graph:



From the graph we can see that when the number of observations are low, we essentially memorize the training data, making the training error almost 0, while having a high testing error. As we increase the number of observations, training error starts to steadily increase, while the testing error drops to a lower plateau. For a large number of observations, the testing error is constant, meaning the number of observations is not a deciding factor in the accuracy of the tree in this case. However, after a certain point, each leaf is making too many observations, generalizing the results too much, causing a large increase in both training and testing error.

2 ORDINARY LEAST SQUARES VERSUS ROBUST LINEAR REGRESSION

When implementing the Ordinary Least Squares method, the input data matrix will in general yield a unique solution. A different input data matrix will return new values of W and b . This occurs because when determining the values of W , they are multiples of the input data matrix.

The value of $w_{OLS} = 1.2476$. The value of $b_{OLS} = 2.528$. The calculated value of $MSE = 1.390425$. The calculated value of $MAD = 0.965729$.

For cauchy, $MSE = 0.299548$, $MAD = 0.242677$.

For fair, $MSE = 0.286079$, $MAD = 0.246835$.


For huber, $MSE = 0.292181$, $MAD = 0.245091$.

For talwar, $MSE = 0.302416$, $MAD = 0.243493$.

Compared to the above values for OLS, we see the MSE and MAD values are much lower for robustfit. This is due to the fact that robustfit ignores outliers and focus on the majority of proximal points. However, OLS takes everything into account, making an outlier have a large effect on the overall curve.

The values for huber are: $w_{huber} = 0.127472$, $b_{huber} = 3.094357$.

We can see the overall graph here:



As mentioned above, we see the 4 robustfit functions are very closely aligned to the majority of points, ignoring outliers. However, OLS is deeply influenced by the outliers, and thus does not make a good fit for the majority of the data.

3 OVERFITTING AND RIDGE REGRESSION

4 LASSO VS RIDGE